


NEURAL NETWORKS (PART 1)

Artificial neuron model



$x_1 \dots x_N$ - inputs

$w_1 \dots w_N$ - weights

b - bias

f - activation function


Activation:

$$a = b + \sum_{i=1}^N x_i w_i$$

Output:

$$z = f(a)$$

Artificial neuron model



$x_1 \dots x_N$ - inputs

$w_1 \dots w_N$ - weights

b - bias

f - activation function

Activation:

$$a = \sum_{i=0}^N x_i w_i$$

$$a = w^T x$$

Output:

$$z = f(a)$$


$$z = f(w^T x)$$

$$w = \begin{bmatrix} b \\ w_1 \\ \vdots \\ w_N \end{bmatrix} \in \mathbb{R}^{N+1}$$

$$x = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_N \end{bmatrix} \in \mathbb{R}^{N+1}$$


Activation functions (most common)

$$f(a) = \begin{cases} 0 & a < 0 \\ 1 & a \geq 0 \end{cases}$$




Heaviside step function
Threshold Logic Unit (TLU)
Perceptron
(classification)

$$f(a) = \begin{cases} 0 & a < 0 \\ a & a \geq 0 \end{cases}$$



RELU function
(nonlinear)

$$f(a) = \frac{1}{1 + e^{-a}}$$




Sigmoid function
Feedforward networks
(nonlinear, classification)

FEEDFORWARD NEURAL NETWORK (MULTILAYER PERCEPTRON)

Feedforward layer architecture

- Neurons are organized in layers (Multilayer Perceptrons)
- Input information is propagated from input neurons towards the output ones



input

$$x \in \mathbb{R}^{N+1}$$

output


$$z^{(2)} \in \mathbb{R}^K$$

weights

$$w^{(1)} \in \mathbb{R}^{(N+1) \times H}$$

$$w^{(2)} \in \mathbb{R}^{(H+1) \times K}$$

Feedforward layer architecture



input $x \in \mathbb{R}^{N+1}$
 output $z^{(2)} \in \mathbb{R}^K$
 weights $w^{(1)} \in \mathbb{R}^{(N+1) \times H}$
 $w^{(2)} \in \mathbb{R}^{(H+1) \times K}$

The network implements the function:

Hidden neuron $z_j^{(1)} = f_j^{(1)} \left(\sum_{i=0}^N w_{ij} x_i \right)$


Output neuron $z_k^{(2)} = f_k^{(2)} \left(\sum_{j=0}^H w_{jk}^{(2)} z_j^{(1)} \right) = f_k^{(2)} \left(\sum_{j=0}^H w_{jk}^{(2)} f_j^{(1)} \left(\sum_{i=0}^N w_{ij}^{(1)} x_i \right) \right)$

Hidden neurons (units)

- Situated in **hidden layers** between the input and the output.
- They allow a network to learn **non-linear** functions and to represent combinations of the input features.

TRAINING (LEARNING) AND TESTING

Learning = minimizing training error (loss)




Learning:
weights are optimized
(iteratively updated) while
the training error is
minimized.

e.g.,
$$E = \frac{1}{2} \sum_{k=1}^K e_k^2 = \frac{1}{2} \sum_{k=1}^K (z_k - y_k)^2$$


Training: Backpropagation algorithm

- For learning (updating of randomly initialized weights) the gradient of the error function is needed.
- The gradient of the error function is calculated by the local exchange of messages in 2 passes:
 - **Forward:** Calculate activations and outputs of all neurons z
 - **Backward:** Calculate errors δ and propagate them back



Testing


- **ONLY Forward pass:** Calculate activations and outputs of all neurons



REGULARIZATION IN NEURAL NETWORKS

Regularization of NN

- How many hidden layers and how many neurons?
 - Fewer – risk of underfitting
 - More – risk of overfitting




Weight decay

- “Weight decay” is an L_2 norm regularization for neural networks.
- The weights of a NN will be an additional term in an Error function:

$$E(\mathbf{w}) = MSE(\mathbf{w}) + \frac{\lambda}{2} ||\mathbf{w}||^2$$

Early stopping

- A form of regularization based on the scheme of model selection
- Steps:
 - The weights are initialized to small values
 - Stop when the error on validation data increases




IMPORTANT DETAILS OF NEURAL NETWORK TRAINING

Important details of training

- If you use another library (than sklearn), or implement everything from scratch, you will need to choose:
 - Activation function also in the *output* layer
 - Loss function appropriate for your data set
 - Batch size - number of training examples used in forward-pass, after which the backward-pass is performed.
- In case of multi-class classification, sometimes the labels are discrete numbers (e.g., if number of classes $n = 10$, labels are discrete numbers 0-9), but sometimes they need to be transformed into “1-out-of- n ” coding scheme.

Regression




- Linear activation function in output layer ensures arbitrary output range

Activation
functions



Classification



- Sigmoid activation function in output layer ensures outputs between 0 and 1

Activation
functions



Classification: binary vs multiclass

- Binary classification:
 - Round the output of a single neuron (with sigmoidal activation) to 0 or 1 (threshold at 0.5) and interpret it as a class 0 or class 1
- Multiclass classification:
 - Multiple output neurons: use 1-out-of-n encoding for the target value $\mathbf{y}^{(i)}$ (one of $y_k^{(i)}$ values is 1, all others are 0)
 - This means there is one output neuron for each class
 - Use a softmax encoding to code the output as probabilities

$$\text{softmax}(z)_k = \frac{e^{z_k}}{\sum_{l=1}^K e^{z_l}}$$

For example if:

$$\mathbf{y}^{(i)} \text{ is one of: } \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix}, \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} \text{ or } \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

We would want the output of ANN to be:

$$\mathbf{z} \approx \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix}, \mathbf{z} \approx \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix}, \mathbf{z} \approx \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

Different error functions for different tasks

- **Regression**

- Quadratic loss:

$$E^{(i)} = \frac{1}{2} \sum_{k=1}^K (z_k^{(i)} - y_k^{(i)})^2$$

- **Binary Classification**

- Binary cross-entropy:

$$E^{(i)} = -y^{(i)} \log(\text{sigmoid}(z^{(i)})) - (1 - y^{(i)}) \log(1 - \text{sigmoid}(z^{(i)}))$$


- **Multi-class Classification**

- Multi-class cross-entropy:

$$E^{(i)} = - \sum_{k=1}^K y_k^{(i)} \log(\text{softmax}(z^{(i)})_k)$$

Training data:

$$\langle \mathbf{x}^{(1)}, \mathbf{y}^{(1)} \rangle \dots \langle \mathbf{x}^{(m)}, \mathbf{y}^{(m)} \rangle$$



Batch vs online learning

- **Batch learning**

- The error gradient for each sample from training set is calculated and accumulated. The weight update is done after all samples are seen.

$$E = \sum_i^m E^{(i)} \quad w_{kj} := w_{kj} - \eta \nabla E$$

- **Online learning**

- After presentation of each sample i from the training set we use the calculated error gradient for weight update:

$$w_{kj} := w_{kj} - \eta \nabla E^{(i)}$$

- It can be used when there is no fixed training set (new data keeps coming in)
- The noise in the gradient can help to escape from **local minimum**
- “Stochastic Gradient Descent”

Mini-batch learning

- Error gradient is calculated and accumulated over samples of a mini-batch from the training set. Weight update is done each time after a **mini-batch** of samples are seen.

$$w_{kj} := w_{kj} - \eta \sum_{i=1}^{N_B} \nabla E^{(i)}$$

- The algorithm is executed (usually in epochs during which a batch of samples from the training set is presented to the network) until a stopping criteria (e.g., error is smaller than some threshold) is satisfied.
- Results may converge to a local minimum.